



Origin LeWM: Encoder features

median  $r/NN$  1.41 ·  $r < NN$  22/128 (17.2%) · fully disjoint 0/128 (0.0%)




Origin LeWM: 8-step rollout predicted features

median  $r/NN$  1.86 ·  $r < NN$  7/128 (5.5%) · fully disjoint 0/128 (0.0%)




Noise-trained LeWM: Encoder features

median  $r/NN$  0.10 ·  $r < NN$  125/128 (97.7%) · fully disjoint 122/128 (95.3%)



Noise-trained LeWM: 8-step rollout predicted features

median  $r/NN$  0.19 ·  $r < NN$  122/128 (95.3%) · fully disjoint 108/128 (84.4%)



NN denotes nearest clean-anchor spacing; all subtitle metrics use the original high-dimensional representation.