

Action-Conditioned Predictive Diagnostics for JEPA World Models

Author names to be supplied for arXiv v1

July 2026

Abstract

Latent predictive world models avoid reconstructing pixels, but visual changes that leave the state unchanged—noise, blur, or reduced resolution—can still alter the encoded history, and the resulting error can grow as the model predicts forward and a planner acts on those predictions. We introduce Action-Conditioned Predictive Consistency (ACPC), a diagnostic for frozen checkpoints. ACPC encodes clean and perturbed views of the same history, predicts both forward under the same actions, and measures the distance between the two latent trajectories. Exact samplewise inequalities relate this distance to changes in multi-step prediction error and planner candidate costs. Because a collapsed representation makes ACPC artificially small, two checkpoint-level summaries complement it: Action-Conditioned Trajectory Radius (ATR, same-state visual sensitivity) and Selective Margin Pass Rate (SMPR, separation of pairs with different logged task states). Across four visual control tasks on the JEPA world model LeWM, eight-step ACPC under the recorded actions predicts perturbation-induced error changes with $51.3 \pm 3.5\%$ lower cross-validated MAE than the strongest of three same-horizon controls with destroyed action information. Candidate-specific ACPC also improves cross-task prediction of CEM selection regret: cross-task test MAE decreases by $15.2 \pm 2.0\%$, with an improvement in all 12 task-run evaluations. A single ATR-SMPR threshold range, recovered from almost any subset of tasks, screens checkpoints across the four tasks to within one augmentation level, orders checkpoint pairs correctly under blur and resize, and applies to PLDM after model-family recalibration. Code is available [here](#).

Keywords: visual world models; predictive stability; action-conditioned diagnostics; cross-task screening; visual corruption.

1 Introduction

World models support control by predicting how candidate actions change the environment [1, 2]. Joint-embedding predictive architectures (JEPAs) instead learn by predicting in representation space rather than reconstructing pixels [3, 4, 5]. This can reduce pressure to model nuisance detail [6, 7]. This is appealing for visual control, where lighting, texture, or image quality may change even when the physical situation does not. Latent prediction alone, however, does not establish that a trained world model will produce stable action-conditioned predictions under such changes.

For control, whether a visual difference is irrelevant depends on the state and the actions being considered. A background texture can be ignored, whereas a contact pose, doorway state, or object-goal relation may determine which action is useful. Encoder invariance cannot resolve this distinction on its own because a planner does not act directly on the encoder output. It rolls the representation forward under candidate actions and ranks the resulting predictions with a cost. A small encoder discrepancy may grow through this computation, whereas a larger one may contract before it affects the cost. This motivates the central question: after the same action sequence is applied, do clean and perturbed views of one trajectory still produce similar predicted futures? A descriptive PushT case study (Section 4.2) shows what is at stake: for a checkpoint trained without noise augmentation, the representations of a history’s perturbed views spread wider than its distance to the nearest other history, both at the encoder and after prediction, so visual noise entangles distinct states.

Action-Conditioned Predictive Consistency (ACPC) turns the central question into a measurement: roll the clean and perturbed views of one history forward under the same actions and measure the distance between the two predicted latent trajectories. Computing the score needs neither an observed future nor a task outcome, yet

two exact inequalities connect it to both: ACPC bounds how much the perturbation can change the multi-step prediction error against a common logged future, and it bounds how far each planner candidate’s cost can move, which yields fixed-pool stability certificates. A small ACPC value alone can still be trivial—a constant representation makes every paired rollout identical—so we summarize each checkpoint with two complementary numbers: Action-Conditioned Trajectory Radius (ATR), a high quantile of the normalized ACPC values over logged evaluation windows, and Selective Margin Pass Rate (SMPR), the fraction of pairs labeled as different states by task coordinates that remain outside that same-state radius by a margin. SMPR certifies only the supplied coordinate labels. A calibrated score combining ATR and SMPR screens frozen checkpoints; ACPC is an evaluation tool, not a training objective.

We instantiate the diagnostic on JEPA-style world models, with LeWM [8] as the primary testbed and PLDM as a second model family (Section 4.7), using Gaussian observation noise for calibration; tasks and protocols are specified in Section 4.1. On LeWM’s unaugmented checkpoints, eight-step ACPC under the recorded actions predicts the perturbation-induced change in prediction error about 51% better, in cross-validated MAE, than the strongest control with the same rollout length but destroyed action information (Section 4.4). ACPC computed along planner candidates consistently improves prediction of CEM selection regret (Section 4.5). Thresholds chosen on some tasks screen checkpoints from tasks excluded from threshold selection, and every partition that selects thresholds on two or more tasks recovers the same threshold pair (Section 4.6); and the calibrated score orders checkpoint pairs correctly under blur and resize at the tested severities (Section 4.8).

The contributions are:

1. **A paired-rollout diagnostic** (Section 3). ACPC measures how same-state visual changes propagate through action-conditioned predictions; ATR summarizes the checkpoint-level radius, and SMPR guards against constant-representation collapse.
2. **Exact links to prediction and planning quantities** (Propositions 1 and 2). Samplewise inequalities connect rollout disagreement to common-future error drift and to squared latent goal-cost movement.
3. **Controlled evidence beyond threshold-selection tasks** (Section 4). Experiments show that ACPC predicts perturbation-induced error changes and CEM selection regret, that a common threshold range recovered from almost any subset of tasks screens checkpoints across tasks, and that the diagnostic can be evaluated under blur, resize, and a second model family.

2 Related Work

Prior work addresses three complementary parts of this problem: learning latent predictive representations, preserving control-relevant distinctions under visual shifts, and evaluating learned dynamics.

Latent predictive world models. DreamerV3 and TD-MPC2 illustrate how learned dynamics can support control by predicting the consequences of candidate actions [1, 2]. JEPAs replace pixel reconstruction with prediction in representation space, first for images and then for video [3, 4, 5, 9]. LeWM instantiates this idea as an end-to-end JEPA world model stabilized by SIGReg [8]; SIGReg’s Gaussianity statistic is based on the Epps–Pulley test [10]. PLDM provides a different joint-embedding dynamics design [11, 12]. Analyses of latent prediction show, in simplified settings, how predictive learning can favor influential features over noisy ones [6, 7]. Under Gaussian, stationary additive-noise dynamics, LeJEPA can recover latent state up to a linear transformation [13]. seq-JEPA studies the related trade-off between invariant and equivariant representations by conditioning sequential prediction on known view transformations, represented as actions [14]. These results motivate latent predictive learning, but our focus is checkpoint evaluation: whether the discrepancy between two visual views of the same trajectory contracts or grows as the predictor rolls them forward under shared actions.

Robust visual control and selective invariance. DrQ, DrQ-v2, and SODA improve visual control through training-time augmentation [15, 16, 17]. ViGMO and VIBR learn invariances in latent or value space [18, 19], whereas ReOI modifies observations at test time to reduce distractor sensitivity in visual model-predictive control [20]. These methods seek to train or repair a more robust controller. We instead ask how to audit visual sensitivity in an already trained predictive model.

Control-aware representation learning also explains why robustness cannot mean removing every difference. Bisimulation-based methods ground behavioral similarity in task-relevant transition structure, often through reward and transition-distribution differences [21, 22, 23, 24], while value-equivalent and value-aware formulations

organize model learning around downstream value or planning quantities [25, 26]. This motivates the two-sided structure of our diagnostic: same-state visual views should agree after prediction, but low disagreement is considered meaningful only when pairs marked as different by task state coordinates remain separated. Our state-coordinate test is narrower than a learned bisimulation or semantic criterion; its role is to expose the constant-representation failure mode.

Diagnostics of learned dynamics. Recent work probes complementary failure modes in learned dynamics. ATM compares action information in encoded and predicted transitions [27]; Delta-JEPA decodes actions from latent differences [28]; ACID tests cycle action consistency with inverse dynamics [29]; Future-Compatible evaluates whether generated futures agree with their claimed actions [30]; and World Models as Group Actions tests identity, inverse, and composition structure [31]. MWM uses observation-level action-conditioned rollout consistency as a training objective [32], while other methods diagnose kinematic failures in long-horizon rollouts [33]. These studies motivate testing the predictor directly instead of inferring its behavior from encoder geometry.

ACPC contributes a complementary paired-view test. It holds the action sequence fixed and measures how much a state-preserving visual change alters the predicted trajectory. We relate this change theoretically to prediction-error drift and fixed-pool candidate-cost movement, and test empirically whether it helps predict adaptive-CEM selection regret. For checkpoint screening, ATR summarizes visual sensitivity and SMPR tests separation on selected task-state coordinates. We calibrate their numerical thresholds separately within each model family.

3 Action-Conditioned Predictive Consistency as a Diagnostic

A visual perturbation can change the pixels without changing the underlying control state. Even so, it can move the encoded history, and the learned dynamics can amplify or contract that change before the planner computes a cost. We study this process at three levels. First, a local geometric analysis compares perturbed versions of one history with other clean histories. Second, ACPC follows a clean-perturbed pair through the complete predicted rollout under the same actions. Third, ATR summarizes ACPC across logged histories, while SMPR checks whether selected pairs with different task-state labels remain separated. ATR and SMPR must be interpreted together: low visual sensitivity is useful only if the representation retains relevant state distinctions.

Throughout, the object under study is a frozen checkpoint—one trained world model with frozen parameters θ . Unlike MWM [32] and seq-JEPA [14], which optimize action-conditioned consistency during training, our paired-view construction is post hoc.

3.1 Local representation geometry under visual perturbations

Before defining the controlled rollout metric, we describe the local geometric picture that motivates it. Start with one clean history and several visually perturbed versions of it; encode or roll forward each version, and compare two scales: the sampled *visual radius*—how far the perturbed representations spread from the clean one—and the *nearest-clean spacing*—the distance from the clean representation to the nearest other clean history. Their ratio, abbreviated r/NN , indicates whether the perturbed views of one history stay closer to it than genuinely different histories do; a symmetric two-ball test additionally counts anchors whose perturbation cloud is fully disjoint from every other anchor’s cloud. Appendix B states the formal definitions and reported summaries; Section 4.2 presents the resulting case study.

These quantities are descriptive. The nearest clean anchor is chosen in learned space and need not lie across a task-relevant boundary; the spacing between clean rollout centers mixes state and action effects, since each center follows its own recorded actions; encoder and endpoint views omit all intermediate predictions; the sampled radius depends on the number of views, anchor density, and representation scale; and a collapsed representation makes every radius small. None of them bounds changes in prediction error or planner cost. They motivate ACPC, which compares clean and perturbed views of the same history under exactly the same action sequence throughout the predicted rollout; when $\alpha_H > 0$, the final-step displacement of any view is at most $\text{ACPC}_H/\sqrt{\alpha_H}$.

3.2 Paired rollout radius

ACPC asks a direct question: if only the visual appearance of a history changes, how much does the model’s predicted future change under a fixed action sequence? Let h be the clean observation–action history consumed by the encoder of the checkpoint under evaluation, and let $\tilde{h}$ be a visually perturbed view of the same underlying state trajectory. The frozen encoder gives $z = E_\theta(h)$ and $\tilde{z} = E_\theta(\tilde{h})$. Both branches receive the same action sequence $\mathbf{a} = (a_0, \dots, a_{H-1})$. For its length- k prefix, define

$$\hat{z}_k = F_\theta^k(z, \mathbf{a}_{0:k-1}), \quad \hat{\tilde{z}}_k = F_\theta^k(\tilde{z}, \mathbf{a}_{0:k-1}). \quad (1)$$

Here F_θ is the action-conditioned predictor and H is the number of autoregressive future steps. The planner does not read raw predictor outputs: it scores candidates in the projected embedding space in which planner costs are computed. The map Π denotes that projection; our implementation stores the projected embeddings directly, so Π requires no extra computation. With nonnegative weights $\sum_{k=1}^H \alpha_k = 1$, stack the whole rollout as

$$\bar{G}_\mathbf{a}(z) = [\sqrt{\alpha_1}\Pi(\hat{z}_1)^\top \quad \dots \quad \sqrt{\alpha_H}\Pi(\hat{z}_H)^\top]^\top. \quad (2)$$

Action-Conditioned Predictive Consistency (ACPC) is the distance between the two weighted predicted rollouts:

$$\text{ACPC}_H(h, \tilde{h}, \mathbf{a}) = \|\bar{G}_\mathbf{a}(E_\theta(h)) - \bar{G}_\mathbf{a}(E_\theta(\tilde{h}))\|_2 = \left(\sum_{k=1}^H \alpha_k \|\Pi(\hat{z}_k) - \Pi(\hat{\tilde{z}}_k)\|_2^2 \right)^{1/2}. \quad (3)$$

Encoder shift $\|z - \tilde{z}\|_2$ measures the effect of the visual change before prediction. ACPC measures its effect on the predicted trajectory after both histories have been propagated under the same actions. Holding the actions fixed isolates sensitivity to the visual perturbation; it does not establish whether the model responds correctly to other actions. Unless stated otherwise we use uniform weights $\alpha_k = 1/H$ and the fixed protocol horizon $H = 8$, the longest horizon for which every logged evaluation window contains a complete observed future (Appendix C). The planner analyses in Section 4.5 instead use $H = 5$, matching the five-action planning horizon of the evaluated CEM configuration.

3.3 Common-future error drift

ACPC can be computed without observing the true future. When a logged future is available, however, the same quantity bounds how much the visual perturbation can change prediction error. Let $Y_\mathbf{a}^H$ be one actually observed future under the same recorded action sequence, projected, stacked, and weighted exactly as in Equation (2). This target is not produced by either prediction branch and is used only for evaluation. Comparing both branches with this one future isolates the error change caused by the visual perturbation.

Proposition 1 (Common-future error-drift bound). *Define*

$$e_h = \|\bar{G}_\mathbf{a}(E_\theta(h)) - Y_\mathbf{a}^H\|_2, \quad e_{\tilde{h}} = \|\bar{G}_\mathbf{a}(E_\theta(\tilde{h})) - Y_\mathbf{a}^H\|_2.$$

Then, for every paired sample,

$$|e_{\tilde{h}} - e_h| \leq \text{ACPC}_H(h, \tilde{h}, \mathbf{a}). \quad (4)$$

The adverse increase $(e_{\tilde{h}} - e_h)_+$, where $(x)_+ = \max(x, 0)$, obeys the same bound.

The proof—a direct application of the reverse triangle inequality in the weighted projected-rollout space, with both errors measured against the same $Y_\mathbf{a}^H$ —is spelled out in Appendix A.

The proposition bounds the difference between two errors, not either prediction’s absolute error: both predictions may be inaccurate even when ACPC is small. The prediction experiment in Section 4.4 uses a logged future only to evaluate whether ACPC is informative about the bounded quantity, which we call the *error drift* $d = |e_{\tilde{h}} - e_h|$.

3.4 Candidate-cost drift and planner stability

We next connect ACPC to planning. A visual perturbation moves each candidate’s predicted endpoint and therefore may change its cost. If no candidate can move far enough to cross the clean cost gap, the selected

candidate remains unchanged. This is the radius–margin argument formalized below. For candidate action sequence $\mathbf{a}^j$, let

$$x_j = \Pi(F_\theta^H(E_\theta(h), \mathbf{a}^j)), \quad \tilde{x}_j = \Pi(F_\theta^H(E_\theta(\tilde{h}), \mathbf{a}^j))$$

be the final clean and perturbed predicted representations. Their displacement $r_j = \|x_j - \tilde{x}_j\|_2$ is the final-step component of the candidate-specific $\text{ACPC}_H(h, \tilde{h}, \mathbf{a}^j)$. If $\alpha_H > 0$, then $r_j \leq \text{ACPC}_H(h, \tilde{h}, \mathbf{a}^j)/\sqrt{\alpha_H}$; the planner experiments in Section 4.5 record r_j directly.

Proposition 2 (Squared goal-cost drift and fixed-pool stability). *Suppose the planner scores candidate j against fixed goal embedding g by $C_j = \|x_j - g\|_2^2$ and $\tilde{C}_j = \|\tilde{x}_j - g\|_2^2$. Define*

$$b_j = r_j(\|x_j - g\|_2 + \|\tilde{x}_j - g\|_2). \quad (5)$$

Then every candidate satisfies $|\tilde{C}_j - C_j| \leq b_j$.

Let w and $\tilde{w}$ be the clean and perturbed winners of the same pool. The clean-reference selection regret obeys

$$0 \leq C_{\tilde{w}} - C_w \leq b_{\tilde{w}} + b_w. \quad (6)$$

The quantity b_j is a candidate-specific upper bound on the movement of the squared goal cost, and the regret inequality bounds the extra clean-history model cost of choosing the perturbed winner; both use the evaluated endpoints, not a global Lipschitz approximation. Subtracting the allowed movements from each clean candidate gap yields selection certificates: if every remainder is positive, the clean winner—and, under the analogous elite-wise condition, the exact top- k elite set—is preserved (Proposition 3). Combining b_j with the final-step conversion $r_j \leq \text{ACPC}_H(h, \tilde{h}, \mathbf{a}^j)/\sqrt{\alpha_H}$ gives the explicit chain

$$|\tilde{C}_j - C_j| \leq b_j \leq \frac{\text{ACPC}_H(h, \tilde{h}, \mathbf{a}^j)}{\sqrt{\alpha_H}}(\|x_j - g\|_2 + \|\tilde{x}_j - g\|_2), \quad (7)$$

so candidate-specific ACPC directly bounds every candidate-cost movement, and through Equations (6), (15) and (16) the planner-facing statements as well. Proofs and equality cases are in Appendix A. Because every candidate must be evaluated under both histories, these certificates support a post-hoc audit rather than a cheap online pre-screen.

These certificates extend to whole planner runs. CEM repeatedly fits its next proposal to the current elite set [34]; with common random numbers, a certified elite set at every iteration aligns the clean and perturbed runs—pools, proposals, and returned action (Corollary 1). The guarantee stops at the first iteration whose elite condition fails, and it concerns candidate selection and model cost, not simulator return. Formal statements and proofs are in Appendix A.

3.5 Checkpoint-level radius and proxy separation

The certificates above are candidate-wise: they require paired evaluation of every candidate in a planner pool. The checkpoint summaries introduced now serve a different purpose—screening—and do not estimate that pool-wise bound. ACPC measures one clean–perturbed history pair; to evaluate a checkpoint, we must summarize many such measurements. A constant representation makes every ACPC value zero while removing all state distinctions, the familiar collapse failure of view-agreement objectives [35]. These two requirements yield a two-part summary: the Action-Conditioned Trajectory Radius (ATR) aggregates sensitivity across histories, and the Selective Margin Pass Rate (SMPR) tests whether selected pairs with different task-state labels remain separated. This two-sided structure parallels bisimulation-based representation learning for control, which grounds behavioral similarity in reward and transition-distribution differences [22].

Evaluation uses a fixed set of *anchors*. Each anchor consists of one logged history window, its H -step observed future, and its recorded actions. We first normalize ACPC separately for each anchor and then aggregate across anchors. Throughout this section, Q_q denotes the empirical q -quantile, n is the number of anchors, and M is the number of perturbation draws per anchor. Distances for anchor i are normalized by its clean-motion scale s_i : the median of its H consecutive observed clean-frame displacements in the projected space (Equation (23), Appendix C). Thus s_i is the anchor’s typical one-step clean motion; dividing by it expresses rollout distances in

units of ordinary motion and uses no task outcome. With a small numerical stabilizer $\varepsilon_{\text{norm}}$ (value in Appendix C) and perturbation draw m , define

$$R_i^{(m)} = \frac{\text{ACPC}_H(h_i, \tilde{h}_i^{(m)}, \mathbf{a}_i)}{s_i + \varepsilon_{\text{norm}}}, \quad \bar{R}_i = \frac{1}{M} \sum_{m=1}^M R_i^{(m)}. \quad (8)$$

The raw *Action-Conditioned Trajectory Radius* (ATR) is

$$\text{ATR}_q^{\text{raw}}(\theta) = Q_q(\{\bar{R}_i\}_{i=1}^n). \quad (9)$$

ATR is one number per checkpoint. It reports a high quantile of normalized ACPC across the sampled histories and therefore emphasizes histories with relatively high visual sensitivity. We use q90 because a mean can hide a small set of high-radius anchors; checkpoint-level conclusions are stable across horizons $H \in \{1, 2, 4, 8\}$ and quantiles q80–q95 (Table 2). This is a descriptive summary, not a candidate-distribution tail guarantee.

ATR alone cannot rule out a collapsed representation. SMPR provides the complementary separation test. Let y_i be the standardized logged task-state vector at the end of anchor i ’s observed H -step future, and let $\ell(y_i)$ be a declared task-specific coordinate label. Each eligible anchor selects its nearest different-label neighbor $j(i)$ inside a fixed state-space neighborhood, both clean histories are rolled under the anchor’s actions $\mathbf{a}_i$, and D_i^{diff} denotes their normalized weighted-rollout distance (Equation (24), Appendix C); the shared metric, action sequence, and normalization put visual sensitivity and state separation on a comparable scale. The neighborhood radius and the margin below are protocol constants fixed a priori. The *Selective Margin Pass Rate* (SMPR) is

$$\text{SMPR}_{q,\delta}(\theta) = \frac{1}{|\mathcal{I}|} \sum_{i \in \mathcal{I}} \mathbf{1}[D_i^{\text{diff}} > \text{ATR}_q^{\text{raw}}(\theta) + \delta], \quad (10)$$

where $\mathcal{I}$ contains anchors with an eligible label-crossing neighbor. The reported protocol fixes $q = 0.90$ and normalized margin $\delta = 0.10$. SMPR is the fraction of tested pairs with different proxy labels whose rollout distance exceeds the same-state radius by more than this margin.

The local ball analysis (Section 3.1) also compares a radius with a spacing, but it is a descriptive case study of learned-space neighborhoods; SMPR instead selects neighbors by a declared task-state label, rolls both histories under the anchor’s actions, and tests complete weighted trajectories against the checkpoint-level q90 same-state radius.

ATR and SMPR must be read together. A constant representation makes both ATR and every D_i^{diff} zero, so it fails the SMPR test. Conversely, a model can separate the tested proxy classes while remaining visually sensitive, which yields a large ATR. These proxies test only the declared coordinate distinctions; they do not prove semantic validity, action relevance, or correct behavior. Appendix C gives the exact coordinate rules and eligible-pair counts.

3.6 Checkpoint-level threshold selection

The raw ATR in Equation (9) is the same-state reference used by SMPR. Raw ATR scales differ across tasks, so comparisons between augmentation levels use a relative value. Let θ_0 be the unaugmented checkpoint for the same task, training run, and model family. Its raw ATR is strictly positive in every evaluated setting, so the following ratio is well defined:

$$\widetilde{\text{ATR}}_q(\theta; \theta_0) = \frac{\text{ATR}_q^{\text{raw}}(\theta)}{\text{ATR}_q^{\text{raw}}(\theta_0)}. \quad (11)$$

Raw ATR remains in the SMPR definition (Equation (10)); relative ATR equals one for the unaugmented reference and removes a common multiplicative scale factor without establishing absolute comparability across tasks or model families.

We combine relative ATR and SMPR only after selecting one threshold for each. For model family $\mathcal{F}$ and visual shift v (e.g., Gaussian noise), let these thresholds be t_R and t_M . Every quantity in the score is measured under the same shift. With a small constant $\varepsilon_{\text{score}}$ (value in Appendix C), define

$$\mathcal{S}_{\mathcal{F},v}(\theta; \theta_0) = \min \left\{ \frac{t_R - \widetilde{\text{ATR}}_q(\theta; \theta_0)}{|t_R| + \varepsilon_{\text{score}}}, \frac{\text{SMPR}_{q,\delta}(\theta) - t_M}{|t_M| + \varepsilon_{\text{score}}} \right\}. \quad (12)$$

The score satisfies $S \geq 0$ exactly when both conditions pass: relative ATR is at most t_R , and SMPR is at least t_M . No task outcome is needed to score a checkpoint once the thresholds are fixed. Planning success rates are used only to select thresholds on threshold-selection tasks and evaluate them on test tasks. Here, *calibration* means operational threshold selection on designated tasks—not probabilistic calibration—and is performed separately for each model family; the equations do not imply that architectures share numerical thresholds.

For a comparison with reference checkpoint θ_0 , define the paired score change $\Delta S_{\mathcal{F},v}(\theta_1; \theta_0) = S_{\mathcal{F},v}(\theta_1; \theta_0) - S_{\mathcal{F},v}(\theta_0; \theta_0)$. A positive ΔS means only that θ_1 improves on its reference under the same calibrated scoring map; it assigns no absolute robustness label to either checkpoint.

Table 1 summarizes the chain. The prediction experiment (Section 4.4) evaluates whether ACPC is informative about the error drift bounded by Proposition 1; the CEM experiment (Section 4.5) asks whether candidate-specific ACPC helps predict adaptive-CEM selection regret; and the cross-task, PLDM, and blur/resize analyses (Sections 4.6 to 4.8) evaluate the calibrated ATR–SMPR score. These are related but distinct empirical claims.

Table 1: Reader guide to the analysis quantities and where each enters the experiments.

Object	Plain-language meaning	Empirical role
Local r/NN	Perturbation-cloud size relative to nearby clean histories; ball separation counts fully disjoint clouds.	Geometry case study (Section 4.2)
ACPC	Distance between clean and perturbed predicted futures under shared actions.	Prediction and planner analyses (Sections 4.4 and 4.5)
b_j	Bound on the movement of candidate j ’s squared goal cost.	Planner cost bound (Section 4.5)
ATR^{raw}	q90 of the normalized same-state rollout radius.	Tube used by SMPR
D_i^{diff}	Rollout distance to a nearby state with a different proxy label.	Proxy-separation test
SMPR	Fraction of proxy-different pairs outside the tube by the margin.	Collapse guard (Section 4.3)
$\widehat{\text{ATR}}, S$	Task-relative radius; S is the smaller of its two calibrated margins.	Sweep and cross-task screening (Sections 4.3 and 4.6)
ΔS	Score change relative to a reference checkpoint.	Blur/resize ordering (Section 4.8)

4 Experiments

We first state the common evaluation protocol. We then begin with a descriptive local-geometry case study that motivates the controlled ACPC comparison, before evaluating checkpoint-level behavior, prediction-error drift, planning consequences, and cross-task screening.

4.1 Evaluation setup

Models, tasks, and planning evaluation. LeWM [8] is the primary model family, and PLDM [11, 12] tests whether the diagnostic applies to a second joint-embedding dynamics architecture. We evaluate four control tasks: TwoRoom navigation, PushT planar manipulation, Reacher arm control, and OGBench-Cube (Cube) 3D manipulation. Following the standard LeWM latent-space MPC evaluation, CEM uses the frozen world model to optimize candidate sequences of five actions (plan horizon 5). We report *planning success rate*, the percentage of evaluation episodes that reach the task goal.

We use *visual perturbation* for one sampled change to a history and *visual shift* for the distribution from which that change is drawn. The main evaluation adds Gaussian noise with $\sigma = 0.08$ to the observations while keeping the goal image clean. Additional evaluations use Gaussian blur ($k = 15$) and resize (scale 0.25).

For each model family and task, the Gaussian-augmentation sweep contains nine checkpoints: one *unaugmented checkpoint*, trained without noise augmentation, and eight checkpoints trained with full-sequence noise $\sigma_{\text{max}} \in \{0.01, \dots, 0.08\}$. Each LeWM task-augmentation setting has three training runs (seeds 3072/3073/3074). Each run is evaluated with seeds 42/43/44 and 100 episodes per evaluation seed. We use *task-run cell* for one task and one training run.

The training run is the unit of replication. The three evaluation seeds measure only within-run variability, so the reported means and dispersions across runs are descriptive rather than population estimates. The PLDM sweep has one run for each of its 36 task-augmentation settings, with the same evaluation seeds and episode count.




Figure 1: Descriptive local-geometry case study motivating the controlled ACPC comparison. Panels (a)–(d) compare a matched pair of PushT checkpoints—the left two unaugmented, the right two trained with Gaussian augmentation at $\sigma_{\max} = 0.08$ —at the encoder output and after eight autoregressive rollout steps. Each color denotes one of 16 highlighted histories: the large black-rimmed marker is its clean anchor, smaller same-color markers are its 18 perturbed views, and the faint ellipse is their 90% t-SNE envelope; gray points show all other histories and views. Under strong contraction, smaller markers and ellipses can lie beneath the clean marker. Each panel’s upper-left summary reports, across all 128 anchors, the median r/NN and the percentages with $r < NN$ and with fully disjoint high-dimensional enclosing balls. The t-SNE coordinates and ellipses are qualitative and enter none of these metrics.

Diagnostic measurements. ATR uses 100 logged anchors, five Gaussian perturbation draws at the evaluation severity $\sigma = 0.08$ (blur and resize diagnostics instead apply the corresponding shift), eight predicted steps, uniform horizon weights, within-anchor draw averaging, and a q90 summary across anchors. SMPR selects pairs within a global q35 radius among standardized endpoint states and tests their separation from the same-state q90 reference with margin 0.10. Appendix C gives the complete normalization and threshold protocol. The raw ATR defines the tube in SMPR; checkpoint thresholding uses the task-relative ATR in Equation (11).

Threshold evaluation. A Gaussian-sweep checkpoint meets the *success-rate criterion* if it recovers at least 80% of the improvement from the unaugmented checkpoint to the best checkpoint at evaluation noise $\sigma = 0.08$, while losing no more than five percentage points on clean observations. Both constants were fixed a priori. We select (t_R, t_M) on every nonempty proper subset of the four tasks and apply the thresholds unchanged to the remaining tasks. The one-, two-, and three-task selection settings produce $4 + 6 + 4 = 14$ directional threshold-selection/test partitions. Metrics first average training runs within each test task and then weight tasks equally; checkpoint rows are never treated as independent replicates. Success rates are used to select and evaluate thresholds, but they are not inputs to ACPC, ATR, or SMPR.

For blur and resize, each task uses the Gaussian thresholds selected on the other three tasks. We compare 24 checkpoint pairs (four tasks $\times$ three runs $\times$ two shifts). Each pair contains the unaugmented checkpoint and the checkpoint trained at $\sigma_{\max} = 0.08$. A pair is positive when success under the shift improves by at least five percentage points and clean success decreases by no more than five points. We make no blur- or resize-specific adjustment. This analysis also does not benchmark alternative signals such as encoder shift or one-step ACPC.

4.2 Motivating case study: local geometry before and after prediction

We begin with the descriptive geometric audit introduced in Section 3.1. Its purpose is to visualize the phenomenon that motivates ACPC, not to establish task-level robustness. Using a fixed PushT case study, we ask whether augmentation makes perturbed views of the same history more compact relative to other clean histories. We compare a matched pair of LeWM checkpoints: one unaugmented and one trained with full-sequence Gaussian augmentation at $\sigma_{\max} = 0.08$, comparing the encoder output with the representation after eight autoregressive rollout steps. All ratios and fractions use the original 192-dimensional projected latent space in which the planner computes costs; the sampling protocol is in Appendix B. The t-SNE visualization [36] is qualitative only.

Figure 1 summarizes the case study. Without augmentation, the median r/NN is 1.41 at the encoder and 1.86 after eight rollout steps, and none of the 128 anchor clouds is fully disjoint. After Gaussian augmentation, the corresponding ratios fall to 0.10 and 0.19; the fully disjoint fractions rise to 95.3% at the encoder and 84.4% after the rollout. Thus, in this fixed case study, augmentation makes perturbed views of the same history substantially




Figure 2: Checkpoint-level planning performance and diagnostics across the complete Gaussian-augmentation sweep. We report planning success rate at evaluation noise $\sigma = 0.08$, relative ATR, and SMPR. Points are across-run means and error bars span the three training runs; success-rate axes are scaled per task. The dashed gray line marks the unaugmented baseline: the success rate of the $\sigma_{\max} = 0$ checkpoint with no perturbation added at training or evaluation time. The $\sigma_{\max} = 0$ endpoint of each curve is this baseline itself; relative ATR equals one there by construction. Lower ATR (relative to the unaugmented checkpoint) and higher SMPR are favorable. Dotted horizontal lines mark the common thresholds $t_R = 0.3$ (ATR) and $t_M = 0.95$ (SMPR) shared across tasks (Section 4.6).

more compact relative to nearby clean histories. Much of this separation remains after prediction.

This case study shows that augmentation can contract perturbation-induced variation relative to nearby clean histories, including after prediction. It does not establish task-semantic separation or planning robustness. We therefore turn to the checkpoint-level ATR–SMPR diagnostic and its relationship with planning performance across the complete augmentation sweep (Figure 2).

4.3 Planning performance under observation noise

Do ATR and SMPR track the recovery of planning performance across the full Gaussian-augmentation sweep? At evaluation noise $\sigma = 0.08$, the unaugmented LeWM checkpoints lose 25.9 ± 2.4 , 74.6 ± 5.6 , 41.0 ± 1.4 , and 22.1 ± 2.8 percentage points in success rate on TwoRoom, PushT, Reacher, and Cube, respectively. Gaussian augmentation recovers performance over a range of training levels whose location differs by task. Figure 2 compares success rate with ATR and SMPR across the complete sweep.

Reacher additionally shows a fixed-protocol training effect: from $\sigma_{\max} = 0.02$ onward, its noisy-evaluation success rate matches or exceeds the unaugmented baseline in Figure 2(c), and the same augmentation lifts the unperturbed score from 59.2% to 76–82%. Part of this crossing therefore reflects a general lift of the augmented checkpoints under the fixed training protocol rather than robustness recovery alone. This reading does not affect the success-rate criterion itself, which is defined relative to the unaugmented and best checkpoints at evaluation noise $\sigma = 0.08$ (Section 4.1), not relative to the crossing of the dashed baseline. We treat the lift as an observation under the fixed training protocol, not a claim about asymptotic training behavior.

Every augmentation level that meets the success-rate criterion has lower ATR and higher SMPR than the corresponding unaugmented checkpoint. Neither diagnostic, however, changes monotonically at every augmentation level. **Takeaway:** recovered planning performance coincides with lower visual sensitivity and greater separation of the tested state pairs. This motivates combining ATR and SMPR. A first-order local-sensitivity estimate of the composed encoder–rollout Jacobian offers one mechanism consistent with the reduced visual sensitivity (Appendix G). The next two experiments separately test whether ACPC uses action information and whether it tracks planning-relevant quantities (Sections 4.4 and 4.5).




Figure 3: Cross-validated MAE for predicting the error drift d on the unaugmented checkpoints; lower is better. Each model is evaluated on trajectory groups excluded from fitting. Small points are training runs, thin lines join the same run, and diamonds are task means. The in-figure legend labels correspond to Base, Base+Control₈, and Base+ACPC₈ as defined in Section 4.4. Base+ACPC₈ is lowest in all 12 task-run cells; latent-error scales are task-specific, so compare within a panel.

4.4 Predicting the perturbation-induced error drift

Does ACPC help predict how much a visual perturbation changes multi-step prediction error? For each history pair, Proposition 1 shows that the two errors against the same logged future can differ by at most ACPC_H . We test whether measured ACPC is informative about this bounded quantity, the error drift $d = |e_{\tilde{h}} - e_h|$. We evaluate at the protocol horizon $H = 8$, the longest horizon with a complete logged common future for every anchor; Table 2 shows that the checkpoint-level conclusions are stable for $H \in \{1, 2, 4, 8\}$.

Data and protocol. The analysis uses the unaugmented checkpoint of each task and training run. Trajectories are partitioned into 16 disjoint groups. Each history pair contributes rows at Gaussian probe severities $\{0.02, 0.05, 0.08\}$ with two perturbation draws; zero-severity probes serve only as identity checks. A ridge model is fitted on 15 groups and evaluated on the remaining group, rotating through all 16, and all rows from one trajectory group stay together. We report the mean absolute error (MAE) on groups excluded from model fitting.

Three nested regressions. Every model contains the probe severity and the encoder-history distance $\|z - \tilde{z}\|_2$.

- **Base** adds one-step ACPC under the recorded action: the information available without a multi-step rollout.
- **Base+Control₈** adds the strongest *destroyed-action control*: eight-step ACPC recomputed with *zeroed* actions (every action set to zero), *swapped* actions (the recorded actions of a different trajectory), or *shuffled* actions (the anchor’s own actions in permuted temporal order). Each control performs the same eight-step rollout, so it carries rollout length without the recorded action sequence. “Strongest” is an oracle choice: in each task-run cell we report whichever of the three controls attains the lowest test MAE, which makes the comparison conservative.
- **Base+ACPC₈** instead adds eight-step ACPC under the recorded actions—the only feature whose action sequence is the one that generated the logged future, and hence the feature matching the right-hand side of Proposition 1.

All eight-step features use uniform weights $\alpha_k = 1/8$ in Equation (3). If Base+ACPC₈ improves on Base, multi-step information helps; if it also improves on Base+Control₈, the gain is attributable to the recorded actions rather than to merely rolling out eight steps.

Results. Base+ACPC₈ attains the lowest cross-validated MAE in all 12 task-run cells (Figure 3). For each training run we average the four task-level relative MAE reductions equally and report the mean $\pm$ sample standard deviation across the three training runs: the reduction is $55.9 \pm 4.7\%$ relative to Base and $51.3 \pm 3.5\%$ relative to Base+Control₈. These results concern prediction of the error drift; they do not compare the absolute forecast accuracy of one-step and eight-step world models. Appendix D reports the task-level values and selected




Figure 4: Adding planner-horizon ACPC improves prediction of adaptive-CEM selection regret on tasks excluded from regression fitting. Each colored line connects the baseline and expanded-model errors for one test task in one training run; the thick black line is the equal-task average across the three runs. The added feature is candidate-specific ACPC q90 at the planner’s five-step prediction horizon. The baseline already contains severity, training condition, one-step ACPC q90, and the clean best-second-best cost gap. MAE is computed on $\log(1 + \text{selection regret})$.

controls (Tables 5 and 6). **Takeaway:** eight-step ACPC carries information about the bounded error change that neither encoder shift, one-step ACPC, nor any same-horizon destroyed-action control provides.

4.5 Predicting CEM selection regret across tasks

A prediction change matters to planning only if it affects the plan that CEM selects. We therefore test one downstream quantity: *clean-reference selection regret*. Let $\mathbf{a}$ and $\tilde{\mathbf{a}}$ be the final plans selected from the clean and perturbed histories, respectively. We measure $[C_h(\tilde{\mathbf{a}}) - C_h(\mathbf{a})]_+$: the extra clean-history model cost incurred by selecting the perturbed-history plan. We refer to this single-decision quantity as *CEM selection regret*. It uses the model’s squared latent goal cost; it is neither cumulative regret nor simulator return.

The paired CEM branches follow Appendix H: same initial proposal, common random numbers, each branch updating from its own elites. For every candidate, we compute one- and five-step ACPC along its action sequence and take the pool q90; five steps match the planner’s prediction horizon. We compare two ridge regressions. The baseline uses perturbation severity, training condition, candidate-level one-step ACPC q90, and the clean gap between the best and second-best candidates. The expanded model adds candidate-level five-step ACPC q90. Within each training run, both regressions are fitted on three tasks and tested on the remaining task, rotating which task is excluded from fitting. We evaluate MAE on $\log(1 + \text{selection regret})$; Appendix H gives the CEM budget and sample counts.

Adding planner-horizon ACPC lowers cross-task test MAE by $15.2 \pm 2.0\%$ (mean $\pm$ sample standard deviation across training runs). The error decreases in all 12 task-run test cases (Figure 4). Since the two regressions differ only by the planner-horizon feature, this gain shows that the planner-horizon rollout carries additional cross-task predictive information. The comparison does not separate action conditioning from horizon length; it tests the incremental value of the configured planner-horizon ACPC feature. **Takeaway:** candidate-specific ACPC helps predict the clean-model cost of a perturbation-induced CEM selection change across tasks. It does not show that ACPC changes the selected action or improves environment performance.

4.6 A common threshold range across tasks

Both the augmentation sweep and the threshold grid are discrete, so we do not ask whether tasks share an exact decision boundary. The useful question is coarser: does one *range* of thresholds screen checkpoints correctly on every evaluated task? Agreement is measured in two ways. A checkpoint is screened correctly when its pass/fail decision matches the success-rate criterion of Section 4.1, and the *recovery-onset mismatch* is the distance, in training-augmentation grid levels, between the first checkpoint the diagnostic accepts and the first checkpoint

that meets that criterion; it asks the diagnostic to locate recovery within a nearby range rather than to reproduce an identical boundary.

Such a range exists, and Figure 2 shows it directly: at the dotted thresholds $t_R = 0.3$ and $t_M = 0.95$, the relative-ATR curve of every task drops below its line, and the SMPR curve rises above its line, within one grid level of the point where planning success recovers. The leave-task-out evaluation over the 14 threshold-selection/test partitions of Section 4.1 confirms that this range is not an artifact of any single task: 13 of the 14 partitions select exactly this pair, and with two or three threshold-selection tasks the diagnostic locates recovery on test tasks within 0.5 grid levels on average and one level at most (balanced accuracy 0.900, precision/recall 0.913/0.953). Calibration on a single task is less reliable (balanced accuracy 0.855 on average, ranging from 0.723 to 0.913): Reacher alone selects $t_R = 0.1$ and produces mismatches of up to five levels on the other tasks. This exception is a tie-break in fit rather than a failure of the shared pair; Appendix E reports every partition and the mechanism.

Two limits keep this claim modest. First, $t_R = 0.3$ is the largest value in the evaluated grid (Appendix C), so the upper edge of the workable range is unresolved; looser thresholds might begin to accept visually sensitive checkpoints. Second, the evaluation concerns the joint ATR–SMPR screen, not a standalone collapse test: SMPR guards only against the constant-representation failure mode on the supplied state-coordinate proxies. **Takeaway:** one threshold range, recovered from almost any subset of tasks, screens checkpoints across the four tasks to within one augmentation level; this supports a shared operating range, not an exact universal boundary or a general collapse detector.

4.7 Portability to other model families

The equations and diagnostics of Section 3 reference only a frozen encoder, an action-conditioned predictor, and a latent goal cost, so nothing in the procedure is specific to LeWM. We test this portability on PLDM, a joint-embedding dynamics model with a different architecture and training recipe, by rerunning the complete Gaussian sweep with the same paired histories, eight-step rollouts, ATR normalization, SMPR guard, success-rate criterion, and threshold-selection procedure (one training run per setting, so run-to-run variability is not estimated). Calibration stays inside the family: for each PLDM test task, thresholds are selected on the other three PLDM tasks and fixed for its nine checkpoints; no LeWM number is reused.

Figure 5 repeats the Figure 2 view on PLDM, and the same signature appears. Within the family, all four leave-one-task-out selections again land on the common pair $(t_R, t_M) = (0.3, 0.95)$: the diagnostic starts accepting checkpoints exactly at the level where planning recovers on PushT and Cube (Cube’s underlying improvement is modest, about six points, but meets the relative criterion), one grid level after it on TwoRoom, and never on Reacher—there the unaugmented checkpoint already reaches 80% noisy success, the sweep never improves it by more than 1.4 points, and relative ATR never falls below t_R , so there is almost no recovery signal to locate. Because PLDM trains a single run per sweep level, an isolated level can reflect single-run instability rather than the augmentation level itself; the diagnostic tracks the one visible dip (Appendix E). Per-task balanced accuracies for both families are tabulated in Table 8 (Appendix E). Applying the LeWM thresholds after within-task ATR normalization yields exactly the same decisions—as expected, since both families select the same pair on this grid, although that agreement does not make thresholds architecture-invariant. **Takeaway:** the theory, the diagnostics, and the calibration procedure carry over unchanged to a different training paradigm; the one failing task offers almost no recovery signal to detect.

4.8 Selective transfer of the diagnostic under blur and resize

Blur ($k = 15$) and resize (scale 0.25), each at one fixed severity, test whether the common range of Section 4.6 survives visual shifts that never entered calibration. For every task, the thresholds selected on the other three Gaussian-noise tasks are exactly the common pair $(t_R, t_M) = (0.3, 0.95)$, so the screen carries over without adjustment. For each of the 24 checkpoint pairs in Section 4.1, we recompute ATR and SMPR under the new shift, insert them into the same calibrated score, and predict improvement when the augmented checkpoint scores higher ($\Delta S > 0$; that is, when the noise-trained checkpoint clears the common thresholds with more margin under the shift than its unaugmented counterpart). The observed label requires a success-rate gain of at least five percentage points with at most a five-point clean loss. No blur or resize outcome enters threshold selection.

The two rows of Figure 6 show the correspondence directly: TwoRoom and Reacher transfer strongly, with success-rate gains of 29–56 points and large positive score changes; Cube shows neither a planning gain nor a score gain; and PushT sits between the two groups on both rows, with mixed success changes and small score




Figure 5: The Figure 2 view repeated on PLDM (one training run per setting), with the same common thresholds drawn as dotted lines. Error bars span the three evaluation seeds of the single run, the dashed gray line marks the unaugmented baseline (the success rate of the $\sigma_{\max} = 0$ checkpoint with no added perturbation), and success-rate axes are scaled per task. On TwoRoom, PushT, and Cube the relative-ATR curve falls below t_R and SMPR stays above t_M where planning recovers; on Reacher the noisy success rate barely changes across the sweep and relative ATR never clears t_R , so no checkpoint is accepted.




Figure 6: Correspondence between planning gains and the diagnostic under blur and resize. Bars average the three training runs for each task and shift; per-pair values are in Table 10. Panel (a): change in planning success rate from the unaugmented to the $\sigma_{\max} = 0.08$ checkpoint under the shift. Panel (b): change ΔS in the calibrated score of Equation (12), i.e., how much more margin the noise-trained checkpoint has inside the common thresholds; positive means the diagnostic favors it. Tasks with clear planning gains (TwoRoom, Reacher) also gain score; Cube gains neither.

changes. The magnitude of the score change tracks the size of the success change (Spearman’s $\rho = 0.835$). Quantitatively, the sign of ΔS matches the five-point success-rate label on 22 of 24 pairs; both exceptions are borderline PushT pairs that improve by exactly four percentage points while the score improves, and per-shift classification metrics are in Appendix F. **Takeaway:** where noise training transfers to blur and resize, the diagnostic improves with it, and where it does not, the diagnostic does not; the common range, fixed on Gaussian noise, preserves this ordering at the tested severities without being an absolute robustness classifier.

5 Discussion and limitations

What the evidence shows. The evidence supports three conclusions. First, eight-step ACPC under recorded actions predicts the error drift bounded by Proposition 1 beyond encoder shift, one-step ACPC, and the strongest same-horizon destroyed-action control (Section 4.4). Because every control also uses eight rollout steps, rollout length alone cannot explain this gain. Second, candidate-level five-step ACPC lowers cross-task test error for CEM selection regret beyond severity, training condition, one-step ACPC, and the clean candidate margin, with

an improvement in all 12 task–run evaluations (Section 4.5). Third, thresholds chosen on some tasks retain useful checkpoint-screening accuracy on the other tasks; under blur and resize the transferred score preserves the observed ordering except in two near-criterion pairs, and the same construction carries to PLDM after model-family recalibration (Sections 4.6 to 4.8).

Scope and limitations. The results have several limits. The planner experiment predicts a clean-model cost under paired evaluations; it does not show that using ACPC changes the planner or improves environment success. The theoretical adaptive-CEM guarantee requires the candidate pools to remain aligned and stops at the first failed elite certificate. SMPR tests only the declared coordinate proxies, so a model can still make stable but incorrect predictions. Three training runs per LeWM task, and one run per PLDM setting, support consistency checks but not precise variability estimates. In addition, blur and resize are each tested at one severity, and the one-source results in Section 4.6 show that calibration depends on source-task composition. Finally, the recovery evidence varies a single intervention—the strength of Gaussian noise augmentation—so the diagnostic–behavior alignment is established along that axis; blur, resize, and PLDM probe transfer, but alignment under interventions other than noise augmentation remains untested.

Future work. Important next steps are direct task outcomes for ACPC-informed planner interventions, multiple severities per visual shift, richer different-state proxies, alignment checks under interventions other than noise augmentation, and additional JEPA-family world models with different objectives and planner interfaces; we plan to extend this coverage in future revisions.

6 Conclusion

ACPC measures how much a same-state visual perturbation changes a world model’s predicted trajectory when the actions are held fixed. Exact inequalities relate this change to prediction error and squared latent goal cost. ATR summarizes the perturbation sensitivity of a checkpoint, while SMPR rejects constant-representation collapse on the tested state-coordinate proxies. Across four tasks, ACPC under recorded actions predicts perturbation-induced error changes beyond one-step and destroyed-action controls and improves cross-task prediction of CEM selection regret. The calibrated ATR–SMPR screen approximately reproduces recovery labels on tasks excluded from threshold selection. The transferred score orders checkpoint pairs correctly under blur and resize, and the same construction extends to PLDM after model-family recalibration. ACPC therefore provides a low-cost audit before deploying a frozen checkpoint behind a planner: using only the model and logged data, it detects visual sensitivity that encoder-level checks can miss.

References

- [1] Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse control tasks through world models. *Nature*, 640(8059):647–653, 2025.
- [2] Nicklas Hansen, Hao Su, and Xiaolong Wang. TD-MPC2: Scalable, robust world models for continuous control. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024.
- [3] Yann LeCun. A path towards autonomous machine intelligence. OpenReview preprint, 2022. Version 0.9.2.
- [4] Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 15619–15629, 2023.
- [5] Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mido Assran, and Nicolas Ballas. Revisiting feature prediction for learning visual representations from video. *Transactions on Machine Learning Research*, 2024.
- [6] Hugues Van Assel, Mark Ibrahim, Tommaso Biancalani, Aviv Regev, and Randall Balestriero. Joint-embedding vs reconstruction: Provable benefits of latent space prediction for self-supervised learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 38, pages 21897–21937, 2025.

- [7] Etai Littwin, Omid Saremi, Madhu Advani, Vimal Thilak, Preetum Nakkiran, Chen Huang, and Joshua Susskind. How JEPA avoids noisy features: The implicit bias of deep linear self distillation networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 37, pages 91300–91336, 2024.
- [8] Lucas Maes, Quentin Le Lidec, Damien Scieur, Yann LeCun, and Randall Balestriero. LeWorldModel: Stable end-to-end joint-embedding predictive architecture from pixels. *arXiv preprint arXiv:2603.19312*, 2026.
- [9] Mido Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, Mojtaba Komeili, Matthew Muckley, Ammar Rizvi, Claire Roberts, Koustuv Sinha, Artem Zholus, Sergio Arnaud, Abha Gejji, Ada Martin, Francois Robert Hogan, Daniel Dugas, Piotr Bojanowski, Vasil Khalidov, Patrick Labatut, Francisco Massa, Marc Szafraniec, Kapil Krishnakumar, Yong Li, Xiaodong Ma, Sarath Chandar, Franziska Meier, Yann LeCun, Michael Rabbat, and Nicolas Ballas. V-JEPA 2: Self-supervised video models enable understanding, prediction and planning. *arXiv preprint arXiv:2506.09985*, 2025.
- [10] T. W. Epps and Lawrence B. Pulley. A test for normality based on the empirical characteristic function. *Biometrika*, 70(3):723–726, 1983.
- [11] Vlad Sobal, Jyothir S V, Siddhartha Jalagam, Nicolas Carion, Kyunghyun Cho, and Yann LeCun. Joint embedding predictive architectures focus on slow features. *arXiv preprint arXiv:2211.10831*, 2022. Self-Supervised Learning: Theory and Practice Workshop at NeurIPS 2022.
- [12] Uladzislau Sobal, Wancong Zhang, Kyunghyun Cho, Randall Balestriero, Tim G. J. Rudner, and Yann LeCun. Learning from reward-free offline data: A case for planning with latent dynamics models. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 38, pages 43905–43941, 2025.
- [13] David Klindt, Yann LeCun, and Randall Balestriero. When does LeJEPA learn a world model? *arXiv preprint arXiv:2605.26379*, 2026.
- [14] Hafez Ghaemi, Eilif B. Muller, and Shahab Bakhtiari. seq-JEPA: Autoregressive predictive learning of invariant-equivariant world models. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 38, pages 32943–32973, 2025.
- [15] Denis Yarats, Ilya Kostrikov, and Rob Fergus. Image augmentation is all you need: Regularizing deep reinforcement learning from pixels. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.
- [16] Denis Yarats, Rob Fergus, Alessandro Lazaric, and Lerrel Pinto. Mastering visual continuous control: Improved data-augmented reinforcement learning. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2022.
- [17] Nicklas Hansen and Xiaolong Wang. Generalization in reinforcement learning by soft data augmentation. In *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, pages 13611–13617, 2021.
- [18] Mingyu Park, Samyeul Noh, Hyun Myung, and Donghwan Lee. Zero-shot visual generalization in model-based reinforcement learning via latent consistency. *OpenReview (ICLR 2026 submission)*, 2025.
- [19] Tom Dupuis, Jaonary Rabarisoa, Quoc-Cuong Pham, and David Filliat. VIBR: Learning view-invariant value functions for robust visual control. In *Proceedings of The 2nd Conference on Lifelong Learning Agents*, volume 232 of *Proceedings of Machine Learning Research*, pages 658–682. PMLR, 2023.
- [20] Yuxin Chen, Jianglan Wei, Chenfeng Xu, Boyi Li, Masayoshi Tomizuka, Andrea Bajcsy, and Ran Tian. Reimagination with test-time observation interventions: Distractor-robust world model predictions for visual model predictive control. In *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, 2026.
- [21] Carles Gelada, Saurabh Kumar, Jacob Buckman, Ofir Nachum, and Marc G. Bellemare. DeepMDP: Learning continuous latent space models for representation learning. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 2170–2179. PMLR, 2019.

- [22] Amy Zhang, Rowan T. McAllister, Roberto Calandra, Yarin Gal, and Sergey Levine. Learning invariant representations for reinforcement learning without reconstruction. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.
- [23] Yutaka Shimizu and Masayoshi Tomizuka. Bisimulation metric for model predictive control. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2025.
- [24] Leonardo F. Toso, Davit Shadunts, Yunyang Lu, Nihal Sharma, Donglin Zhan, Nam H. Nguyen, and James Anderson. Learning invariant visual representations for planning with joint-embedding predictive world models. *arXiv preprint arXiv:2602.18639*, 2026.
- [25] Christopher Grimm, André Barreto, Satinder Singh, and David Silver. The value equivalence principle for model-based reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pages 5541–5552, 2020.
- [26] Claas A. Voelcker, Anastasiia Pedan, Arash Ahmadian, Romina Abachi, Igor Gilitschenski, and Amir-Massoud Farahmand. Calibrated value-aware model learning with probabilistic environment models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 61745–61768. PMLR, 2025.
- [27] Jiaheng Chen. ATM: Action-consistency transfer matrix for diagnosing and improving latent world models. *arXiv preprint arXiv:2606.09028*, 2026.
- [28] Zhenghao Zhang, Yuanxiang Wang, Zhenyu Guan, Yujia Yang, Bingkang Shi, Tianyu Zong, Hongzhu Yi, Guoqing Chao, Xingchen Chen, Tiankun Yang, Chenxi Bao, Tao Yu, Jingjing Zhou, and Jungang Xu. Delta-JEPA: Learning action-sensitive world models via latent difference decoding. *arXiv preprint arXiv:2606.31232*, 2026.
- [29] Gawon Seo, Dongwon Kim, and Suha Kwak. ACID: Action consistency via inverse dynamics for planning with world models. *arXiv preprint arXiv:2607.02403*, 2026.
- [30] Bo-Kai Ruan, Teng-Fang Hsiao, Ling Lo, and Hong-Han Shuai. Is the future compatible? Diagnosing dynamic consistency in world action models. *arXiv preprint arXiv:2605.07514*, 2026.
- [31] Zijie Wang, Wei Zhang, Weiming Zhang, Fanqi Zhang, Xiao Tan, Yipeng Qin, and Guanbin Li. World models as group actions. *arXiv preprint arXiv:2605.24578*, 2026.
- [32] Han Yan, Zishang Xiang, Zeyu Zhang, and Hao Tang. MWM: Mobile world models for action-conditioned consistent prediction. *arXiv preprint arXiv:2603.07799*, 2026.
- [33] Finn Rasmus Schäfer, Korbinian Moller, Yuan Gao, Christian Oefinger, Sebastian Schmidt, and Johannes Betz. Imagined rollouts are kinematic, not dynamic: A diagnosis of long-horizon world-model failure. *arXiv preprint arXiv:2607.05966*, 2026. Workshop paper accepted at the RSS Robot World Model Workshop 2026.
- [34] Dirk P. Kroese, Sergey Porotsky, and Reuven Y. Rubinstein. The cross-entropy method for continuous multi-extremal optimization. *Methodology and Computing in Applied Probability*, 8(3):383–407, 2006.
- [35] Adrien Bardes, Jean Ponce, and Yann LeCun. VICReg: Variance-invariance-covariance regularization for self-supervised learning. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2022.
- [36] Laurens van der Maaten and Geoffrey Hinton. Visualizing data using t-SNE. *Journal of Machine Learning Research*, 9(86):2579–2605, 2008.
- [37] Salah Rifai, Pascal Vincent, Xavier Muller, Xavier Glorot, and Yoshua Bengio. Contractive auto-encoders: Explicit invariance during feature extraction. In Lise Getoor and Tobias Scheffer, editors, *Proceedings of the 28th International Conference on Machine Learning*, pages 833–840. Omnipress, 2011.
- [38] Michael F. Hutchinson. A stochastic estimator of the trace of the influence matrix for laplacian smoothing splines. *Communications in Statistics - Simulation and Computation*, 18(3):1059–1076, 1989.

A Proofs and Fixed-Pool Analysis

Proof of Proposition 1. Write $u = \bar{G}_{\mathbf{a}}(E_{\theta}(h))$, $v = \bar{G}_{\mathbf{a}}(E_{\theta}(\tilde{h}))$, and $y = Y_{\mathbf{a}}^H$. All three vectors live in the weighted projected-rollout space of Equation (2), and both errors are measured against the same y . The reverse triangle inequality gives

$$|e_{\tilde{h}} - e_h| = \left| \|v - y\|_2 - \|u - y\|_2 \right| \quad (13)$$

$$\leq \|(v - y) - (u - y)\|_2 = \|v - u\|_2 = \text{ACPC}_H(h, \tilde{h}, \mathbf{a}). \quad (14)$$

Equality holds exactly when $u - y$ and $v - y$ are positively colinear. The adverse increase obeys the same bound because $(e_{\tilde{h}} - e_h)_+ \leq |e_{\tilde{h}} - e_h|$.

Proposition 3 (Fixed-pool selection certificates). *In the setting of Proposition 2, if the clean winner w is unique and*

$$\min_{j \neq w} (C_j - C_w - b_j - b_w) > 0, \quad (15)$$

then the perturbed branch has the same winner. If $\mathcal{E}$ is the clean top- k elite set and

$$\min_{i \in \mathcal{E}, j \notin \mathcal{E}} (C_j - C_i - b_j - b_i) > 0, \quad (16)$$

then the perturbed branch has exactly the same elite set.

Corollary 1 (Conditional CEM stability). *Consider clean and perturbed CEM runs initialized by the same proposal and sampled with common random numbers. While their ordered candidate pools remain aligned, if Equation (16) holds at the current iteration, both runs fit the same next proposal. If it holds at every iteration, their pools and proposals remain aligned by induction; since the evaluated solver returns the final proposal mean as its action, the selected actions then coincide as well.*

Proof of Proposition 2. For one candidate, suppress the index and factor the squared-distance difference:

$$|\tilde{C} - C| = \left| \|\tilde{x} - g\|_2^2 - \|x - g\|_2^2 \right| \quad (17)$$

$$= \left| \langle \tilde{x} - x, \tilde{x} + x - 2g \rangle \right| \quad (18)$$

$$\leq \|\tilde{x} - x\|_2 \|\tilde{x} + x - 2g\|_2 \quad (19)$$

$$\leq r(\|x - g\|_2 + \|\tilde{x} - g\|_2) = b. \quad (20)$$

The first inequality is Cauchy-Schwarz and the second is the triangle inequality. Thus, for nominal winner w and competitor j ,

$$\tilde{C}_j - \tilde{C}_w \geq C_j - C_w - |\tilde{C}_j - C_j| - |\tilde{C}_w - C_w| \geq C_j - C_w - b_j - b_w.$$

For the perturbed winner $\tilde{w}$,

$$C_{\tilde{w}} \leq \tilde{C}_{\tilde{w}} + b_{\tilde{w}} \leq \tilde{C}_w + b_{\tilde{w}} \leq C_w + b_w + b_{\tilde{w}},$$

which proves Equation (6); nonnegativity follows from the optimality of w on the nominal pool.

For Proposition 3, positivity of Equation (15) makes every perturbed competitor strictly more costly than w . The same argument for every $i \in \mathcal{E}$ and $j \notin \mathcal{E}$ proves that Equation (16) preserves the exact elite set. Strict inequalities avoid relying on implementation-specific tie breaking.

Proof of Corollary 1. Index the branches by $b \in \{\text{c}, \text{p}\}$ (clean, perturbed) and iterations by t . Let ϕ_t^b collect branch b 's proposal parameters (mean and variance), let ω_t be the shared random numbers, and let $\mathcal{A}_t^b = T(\phi_t^b, \omega_t)$ be the sampled candidate pool, where T is the deterministic sampling map. Branch b selects the elite set $\mathcal{E}_t^b$ as the k lowest-cost candidates under its own costs and updates $\phi_{t+1}^b = U(\{\mathbf{a}^j : j \in \mathcal{E}_t^b\})$, where U is deterministic and permutation invariant; the evaluated solver's unweighted elite mean and standard deviation satisfy this. The induction step is

$$\phi_t^{\text{c}} = \phi_t^{\text{p}} \implies \mathcal{A}_t^{\text{c}} = \mathcal{A}_t^{\text{p}} \xrightarrow{\text{Equation (16)}} \mathcal{E}_t^{\text{c}} = \mathcal{E}_t^{\text{p}} \implies \phi_{t+1}^{\text{c}} = \phi_{t+1}^{\text{p}},$$

and the base case $\phi_0^c = \phi_0^p$ holds by assumption, which aligns pools and proposals for as long as the certificate holds at every completed iteration. The evaluated solver returns the final proposal mean as its action, so aligned proposals return identical actions. An implementation that instead returns the best final candidate additionally requires the top-1 condition Equation (15) at the last iteration. If a certificate fails, equality of the elite sets is unknown; the induction stops and no later shared-pool statement is made.

Relation to the weighted rollout radius. The planner bound uses the displacement at the cost readout horizon. For the weighted stack in Equation (3), $\sqrt{\alpha_H} r_j \leq \text{ACPC}_H$ whenever $\alpha_H > 0$. The planner experiments record r_j directly, avoiding the looser $1/\sqrt{\alpha_H}$ conversion. This does not require a future target, but it does require evaluating the candidate under both matched histories.

Sharpness without additional assumptions. Neither principal bound admits a uniformly smaller right-hand side from the available quantities alone. For a unit vector u , common target $Y = 0$, and two stacked predictions au and bu with $a, b \geq 0$, the reverse-triangle bound in Equation (4) holds with equality. Likewise, setting $g = 0$, $x = au$, and $\tilde{x} = bu$ makes $|\tilde{C} - C| = |b - a|(a + b) = r(\|x - g\|_2 + \|\tilde{x} - g\|_2)$.

B Local-Geometry Case-Study Definitions

This section states the descriptive quantities summarized in Section 3.1 and reported by the case study in Section 4.2. The case study uses 128 logged PushT histories; for each, one clean view and 18 perturbed views are generated—six draws at each probe standard deviation in $\{0.01, 0.04, 0.08\}$. The t-SNE display reproduces the deterministic per-panel projections of the original Appendix visualization.

Let $\mathcal{V} \in \{\text{enc}, \text{roll}\}$ denote either the encoder output or the representation after eight autoregressive rollout steps (the protocol horizon $H=8$; Section 3.2). For anchor history i , let $v_{i,\mathcal{V}}^{(0)}$ be its clean representation and $v_{i,\mathcal{V}}^{(m)}$, $m = 1, \dots, M$, the representations of M perturbed views. In the rollout space, the clean and perturbed views of each anchor use the same recorded action sequence. We define the sampled visual radius, nearest-clean spacing, and their ratio as

$$\begin{aligned} r_{i,\mathcal{V}}^{\text{vis}} &= \max_{1 \leq m \leq M} \left\| v_{i,\mathcal{V}}^{(m)} - v_{i,\mathcal{V}}^{(0)} \right\|_2, \\ d_{i,\mathcal{V}}^{\text{NN}} &= \min_{j \neq i} \left\| v_{i,\mathcal{V}}^{(0)} - v_{j,\mathcal{V}}^{(0)} \right\|_2, \quad \rho_{i,\mathcal{V}}^{\text{local}} = \frac{r_{i,\mathcal{V}}^{\text{vis}}}{d_{i,\mathcal{V}}^{\text{NN}}}. \end{aligned} \quad (21)$$

The figures abbreviate $\rho_{i,\mathcal{V}}^{\text{local}}$ as r/NN . When this ratio is below one, every sampled perturbed view of anchor i lies closer to its clean representation than that representation lies to the nearest other clean anchor. A ratio of at least one means that the largest sampled displacement reaches or exceeds this nearest-clean distance. This is a geometric comparison; the nearest clean anchor need not lie across a semantic or task-relevant boundary.

The ratio uses only anchor i ’s perturbation radius. It does not account for the perturbation radius around the neighboring anchor. We therefore add a symmetric test: the enclosing ball around anchor i is *fully disjoint* if it is separated from the enclosing ball around every other anchor:

$$\left\| v_{i,\mathcal{V}}^{(0)} - v_{j,\mathcal{V}}^{(0)} \right\|_2 > r_{i,\mathcal{V}}^{\text{vis}} + r_{j,\mathcal{V}}^{\text{vis}} \quad \text{for every } j \neq i. \quad (22)$$

Across the sampled anchors, we report three summaries: the median r/NN ratio; the fraction satisfying $r_{i,\mathcal{V}}^{\text{vis}} < d_{i,\mathcal{V}}^{\text{NN}}$; and the fraction satisfying the fully-disjoint condition in Equation (22). The second is a one-sided nearest-center test, whereas the third is a symmetric certificate for the enclosing balls in the original high-dimensional representation. It is unrelated to whether the qualitative t-SNE display ellipses overlap.

C Evaluation and Diagnostic Protocol

This appendix specifies the fixed evaluation and diagnostic protocol used in the main tables.

Evaluation grid. LeWM is evaluated on PushT, TwoRoom, Reacher, and Cube with evaluation seeds 42/43/44 and 100 episodes per seed. The main evaluation perturbs observation pixels with Gaussian noise at standard deviation 0.08 while keeping the goal image clean. Checkpoints trained with noise use full-sequence input-side Gaussian augmentation with $\sigma_{\max} \in \{0.01, \dots, 0.08\}$.

PLDM instantiation. PLDM uses the same four tasks, nine-checkpoint Gaussian grid, evaluation seeds, trajectory count, paired-history construction, $H = 8$ rollout statistic, and SMPR definition. For cross-task analysis, we divide ATR by the value of the corresponding unaugmented PLDM checkpoint. We then select thresholds on the other three PLDM tasks. Success rates from the evaluation task never enter threshold selection. A second test applies the corresponding LeWM threshold pair after the same within-task PLDM normalization. We do not assume that model families share raw thresholds.

Success-rate criterion. Within each task and training run, let P_{base} and P_{best} be the success rates at evaluation noise $\sigma = 0.08$ for the unaugmented checkpoint and the best checkpoint in the sweep. A checkpoint meets the success-rate criterion when its noisy-observation success rate is at least $P_{\text{base}} + 0.8(P_{\text{best}} - P_{\text{base}})$ and its clean success rate is no more than five percentage points below that of the unaugmented checkpoint. Statements about recovered levels in Figure 2 require this criterion to hold in a majority of runs.

ATR protocol. We compute ATR separately for each task, training run, and checkpoint. For anchor i , let $u_{i,0}^{\text{obs}}$ be the embedding of the last clean history observation in the projected space selected by Π , and $u_{i,1:H}^{\text{obs}}$ the embeddings of the next H actually observed clean frames in that space; the clean-motion scale is

$$s_i = Q_{0.50} \left(\left\{ \|u_{i,k}^{\text{obs}} - u_{i,k-1}^{\text{obs}}\|_2 \right\}_{k=1}^H \right). \quad (23)$$

Each anchor contributes one weighted radius over the rollout horizon. Perturbation draws use Gaussian observation noise at $\sigma = 0.08$, matching the behavioral evaluation severity; under blur or resize, draws apply that shift instead. The default uses $H = 8$ and uniform weights $\alpha_k = 1/H$. We divide each radius by the anchor’s q50 clean observed-transition scale, including the transition from the final history observation to the first future observation. We then average multiple noise draws within each anchor and take q90 across anchors. This gives the raw ATR in Equation (9), which SMPR uses as its same-state reference. For sweep plots and within-family cross-task calibration, we divide raw ATR by the corresponding unaugmented value as in Equation (11) before computing thresholds or run summaries. The numerical stabilizers are $\varepsilon_{\text{norm}} = 10^{-8}$ in Equation (8) and $\varepsilon_{\text{score}} = 10^{-12}$ in Equation (12). Table 2 reports how the checkpoint-level ATR reduction depends on the rollout horizon and the summary quantile.

Table 2: Horizon and quantile sensitivity of the ATR reduction at fixed evaluation noise $\sigma = 0.08$. Each entry is the median, over the three training runs, of the ratio between the ATR of the checkpoint trained at the highest augmentation level ($\sigma_{\max} = 0.08$) and that of the unaugmented checkpoint; lower means a larger reduction. These values are descriptive and do not retune any threshold.

Task	$q = 0.90$ by horizon				$H = 8$ by quantile		
	H1	H2	H4	H8	q80	q90	q95
TwoRoom	0.05	0.04	0.05	0.06	0.06	0.06	0.06
PushT	0.06	0.07	0.06	0.09	0.07	0.09	0.09
Reacher	0.02	0.03	0.03	0.02	0.02	0.02	0.02
Cube	0.04	0.03	0.04	0.04	0.03	0.04	0.04

SMPR protocol. SMPR follows Equation (10). Let $d_{0.35}$ be the q35 of all off-diagonal Euclidean distances among the standardized endpoint states; this radius, well below the median pairwise distance, keeps every selected pair genuinely nearby in state space while leaving most anchors an eligible different-label neighbor (Table 3 reports the counts). Each eligible anchor chooses the nearest neighbor $j(i)$ with a different label and distance at most $d_{0.35}$. The predictor rolls both clean histories under the anchor actions $\mathbf{a}_i$; for the neighbor these are not the actions it actually executed, but reusing the anchor’s actions keeps a single shared action sequence on both sides. Their normalized proxy-different distance is

$$D_i^{\text{diff}} = \frac{\|\bar{G}_{\mathbf{a}_i}(E_\theta(h_i)) - \bar{G}_{\mathbf{a}_i}(E_\theta(h_{j(i)}))\|_2}{s_i + \varepsilon_{\text{norm}}}. \quad (24)$$

Table 3: Implemented state-coordinate proxies for SMPR. The state is taken at the eighth observed future step and standardized coordinate-wise with full-dataset statistics. Counts give eligible and skipped anchors in the fixed 100-anchor audit; these labels are not semantic or action-relevance annotations.

Task	Logged state field (dim.)	Implemented proxy label $\ell(y)$	Eligible / skipped
TwoRoom	<code>pos_agent</code> (2)	One median bit from the standardized agent x coordinate (slice <code>pos_agent[0:1]</code>).	61 / 39
PushT	<code>state</code> (7)	Three median bits from standardized block x , block y , and block angle (slice <code>state[2:5]</code>).	98 / 2
Reacher	<code>observation</code> (6)	Three median bits from the two standardized joint-velocity coordinates and their Euclidean norm (slice <code>observation[4:6]</code>).	100 / 0
Cube	<code>observation</code> (28)	Three median bits from the first two coordinates and the Euclidean norm of $r = \bar{o}_{0:3} - \bar{o}_{25:28}$, where $\bar{o}$ is the standardized observation.	100 / 0

Table 4: Summary of the Gaussian-augmentation sweep (nine checkpoints per task: no augmentation plus eight noise levels; Figure 2 shows every level). “Best” is the level with the highest mean planning success at evaluation noise $\sigma = 0.08$; arrows give the change from the unaugmented checkpoint to that level. Relative ATR is lower-is-better, SMPR higher-is-better; the last column lists levels meeting the success-rate criterion (Section 4.1).

Task	No-aug. success (%)	Best success (%)	Best $\sigma_{\max}$	Relative ATR	SMPR	Levels meeting criterion
TwoRoom	68.8	97.1	0.08	1.00→0.07	0.11→0.98	0.01–0.08
PushT	7.2	86.8	0.06	1.00→0.11	0.28→1.00	0.03–0.08
Reacher	18.2	83.3	0.07	1.00→0.03	0.37→0.99	0.03–0.08
Cube	43.1	66.0	0.03	1.00→0.26	0.07→0.98	0.03–0.07

The pair passes only if this distance is strictly greater than raw q90 ATR plus 0.10.

Table 3 makes the pairing rule explicit. The logged state is taken at the end of the observed eight-step future and standardized coordinate-wise by the dataset preprocessor. The neighborhood radius is the $d_{0.35}$ defined above, computed in that task’s full standardized state vector. Proxy labels use median splits computed within the fixed 100-endpoint anchor batch (anchor seed 9101); each label therefore describes the endpoint reached by that trajectory’s own recorded actions. After selecting a neighbor, we roll both starting histories under the anchor’s actions for the latent-distance test. We skip an anchor only when no state with a different label lies inside the neighborhood. Pairing depends only on the logged states, so the eligible counts remain constant across checkpoints and training runs.

Threshold grids. For cross-task evaluation, candidate thresholds are $t_R \in \{0.05, 0.075, 0.10, 0.15, 0.20, 0.30\}$ for task-relative ATR and $t_M \in \{0.80, 0.85, 0.90, 0.95\}$ for SMPR. Within each threshold-selection task subset, the selection criterion lexicographically minimizes mean absolute recovery-onset mismatch, false-early and false-late counts, then maximizes balanced accuracy. The selected pair is applied without modification to every test task. No test-task label or blur/resize outcome enters threshold selection.

Reproducibility. The scripts documented in `paper1/scripts/README.md` rebuild all summary figures and tables deterministically. Checkpoint-level diagnostics also require the released checkpoints and datasets. Machine-readable summaries record training seeds, evaluation seeds, and protocol hashes.

D Prediction-Error Scope and Controls

The common-future inequality in Proposition 1 has zero numerical violations in any logged eight-step or candidate-specific five-step task×run×training-condition cell; this is an implementation invariant, not an independent statistical success count.

Logged-future prediction and candidate planning are different estimands: the former predicts the error drift against an independent action-matched observed future, whereas the latter compares candidates from the actual planner pool. We do not infer one result from the other; Section 4.5 supplies the planner-facing evidence with same-pool, candidate-specific features.

The prediction-error analysis is restricted to the unaugmented checkpoints and the listed visual probes. Table 5 reports the per-run relative reductions, and Table 6 the underlying out-of-group MAE values with the selected controls. Every training run improves on all four tasks.

Table 5: Relative reduction in out-of-group MAE from Base+ACPC₈ when predicting the error drift d on the unaugmented checkpoints (protocol and model definitions in Section 4.4). Base+Control₈ uses the per-cell oracle control; higher is better. The last column counts task-run cells in which Base+ACPC₈ is best.

Training run (seed)	vs. Base	vs. Base+Control ₈	Task-run cells improved
3072	55.5%	51.4%	4/4
3073	60.8%	54.8%	4/4
3074	51.4%	47.7%	4/4
mean $\pm$ SD	55.9 $\pm$ 4.7%	51.3 $\pm$ 3.5%	12/12

Table 6: Out-of-group MAE (16-fold leave-one-trajectory-group-out) for predicting the error drift d on the unaugmented checkpoints; lower is better, model definitions in Section 4.4. “Selected control” is the destroyed-action control chosen by the per-cell oracle; “paired wins” counts the folds (of 16) in which Base+ACPC₈ beats Base+Control₈.

Task	run (seed)	Base	Base +Control ₈	Base +ACPC ₈	selected control	paired wins
TwoRoom	3072	2.535	2.099	0.755	shuffled actions	15/16
TwoRoom	3073	2.336	2.015	0.866	shuffled actions	15/16
TwoRoom	3074	2.075	1.911	1.006	swapped actions	13/16
PushT	3072	2.068	2.068	1.762	shuffled actions	11/16
PushT	3073	1.969	2.008	1.315	zeroed actions	13/16
PushT	3074	1.909	1.833	1.439	zeroed actions	13/16
Reacher	3072	1.358	1.097	0.288	shuffled actions	16/16
Reacher	3073	1.399	0.793	0.247	shuffled actions	16/16
Reacher	3074	1.409	1.136	0.263	shuffled actions	16/16
Cube	3072	1.171	1.037	0.489	zeroed actions	14/16
Cube	3073	1.416	1.212	0.500	zeroed actions	14/16
Cube	3074	1.061	1.008	0.552	shuffled actions	14/16

E Cross-Task Threshold Partitions

The 14 rows below are exhaustive: selecting one, two, or three of four tasks as sources creates $4 + 6 + 4$ directional partitions. Each evaluation task retains every training run and all nine checkpoints. The apparent sample size is therefore not inflated by treating checkpoint rows as independent observations.

The single-source Reacher exception in Section 4.6 is a tie-break in fit: under $(0.3, 0.95)$ Reacher’s own onset mismatch is also at most one level, but calibrating on Reacher alone prefers the tighter $t_R = 0.1$ that fits it exactly and does not carry over. Recovered Reacher checkpoints reach relative ATR well below 0.1, while the earliest recovered levels on the other tasks lie above it, so the tighter threshold rejects them and delays the accepted onset.

On the PLDM sweep, the most visible single-run artifact is Reacher at $\sigma_{\max} = 0.05$, where the clean and the noisy score drop together. The diagnostic tracks this dip: the two levels with the lowest clean scores ($\sigma_{\max} = 0.01$ and 0.05) are exactly the two with the highest relative ATR in Figure 5(c), consistent with relative ATR responding to checkpoint quality itself rather than only to the augmentation level.

For PLDM, Table 8 lists the per-task balanced accuracies of the leave-task-out screen of Section 4.7 next to the LeWM values; pooled over all 36 PLDM decisions, balanced accuracy is 0.836 with precision 0.789 and recall 0.882.

F Cross-Stressor Pairs and Discordant Cases

Both discordant rows lie near the success-rate criterion: PushT run 3072 under resize and PushT run 3074 under blur each improve success rate by four percentage points, one point below the five-point criterion, while their selective scores increase. The complete table shows these cases rather than only the aggregate classification metrics.

G Local Gaussian Sensitivity as a Mechanistic Interpretation

For a small isotropic observation perturbation $\delta x \sim \mathcal{N}(0, \sigma^2 I)$, first-order expansion of encoder E followed by rollout map G gives $\Delta G = J_G J_E \delta x + o(\|\delta x\|_2)$ with Jacobians evaluated at the clean input. Writing $A = J_G J_E$

Table 7: Checkpoint screening with thresholds chosen on other tasks. Thresholds are selected on the threshold-selection tasks and applied unchanged to all test tasks. Metrics first average training runs within each test task and then weight tasks equally. Recovery-onset mismatch is measured in training-augmentation grid levels (one level is 0.01 in $\sigma_{\max}$).

Selection tasks	Test tasks	t_R	t_M	BA	P / R	Onset mismatch mean / max
TwoRoom	PushT, Reacher, Cube	0.3	0.95	0.888	0.884 / 0.981	0.3 / 1.0
PushT	TwoRoom, Reacher, Cube	0.3	0.95	0.894	0.902 / 0.956	0.6 / 1.0
Reacher	TwoRoom, PushT, Cube	0.1	0.95	0.723	0.906 / 0.540	2.9 / 5.0
Cube	TwoRoom, PushT, Reacher	0.3	0.95	0.913	0.950 / 0.938	0.6 / 1.0
TwoRoom, PushT	Reacher, Cube	0.3	0.95	0.874	0.853 / 1.000	0.3 / 1.0
TwoRoom, Reacher	PushT, Cube	0.3	0.95	0.887	0.873 / 0.972	0.3 / 1.0
TwoRoom, Cube	PushT, Reacher	0.3	0.95	0.903	0.925 / 0.972	0.3 / 1.0
PushT, Reacher	TwoRoom, Cube	0.3	0.95	0.896	0.901 / 0.935	0.7 / 1.0
PushT, Cube	TwoRoom, Reacher	0.3	0.95	0.912	0.952 / 0.935	0.7 / 1.0
Reacher, Cube	TwoRoom, PushT	0.3	0.95	0.926	0.972 / 0.907	0.7 / 1.0
TwoRoom, PushT, Reacher	Cube	0.3	0.95	0.858	0.802 / 1.000	0.3 / 1.0
TwoRoom, PushT, Cube	Reacher	0.3	0.95	0.889	0.905 / 1.000	0.3 / 1.0
TwoRoom, Reacher, Cube	PushT	0.3	0.95	0.917	0.944 / 0.944	0.3 / 1.0
PushT, Reacher, Cube	TwoRoom	0.3	0.95	0.935	1.000 / 0.869	1.0 / 1.0

Table 8: The same leave-task-out screen in two model families: for each evaluation task, thresholds are selected on that family’s other three tasks and applied unchanged. Chance is 0.5; LeWM values average three training runs, PLDM has one run per setting. Applying score-aligned LeWM thresholds to PLDM after within-task ATR normalization yields identical PLDM decisions.

Evaluation task	Balanced accuracy	
	LeWM	PLDM
TwoRoom	0.935	0.938
PushT	0.917	0.750
Reacher	0.889	0.500
Cube	0.858	0.875

and taking expectations of the leading term,

$$\mathbb{E}[\|A \delta x\|_2^2] = \text{tr}(A \mathbb{E}[\delta x \delta x^\top] A^\top) = \sigma^2 \text{tr}(A A^\top) = \sigma^2 \|A\|_F^2, \quad (25)$$

so $\mathbb{E}[\|\Delta G\|_2^2] \approx \sigma^2 \|J_G J_E\|_F^2$ to first order as $\sigma \rightarrow 0$. Jacobian Frobenius norms have previously been used to quantify and regularize local representation sensitivity and contraction [37]. We estimate $\text{tr}[(J_G J_E)^\top (J_G J_E)] = \|J_G J_E\|_F^2$ with Rademacher probes using the Hutchinson estimator [38]; JVPs avoid materializing the full Jacobian. This explains one mechanism by which Gaussian noise training can contract ACPC: it reduces the composed local sensitivity of the encoder–predictor path (Figure 7). It is a local approximation, not a global robustness bound.

H Adaptive-CEM Selection-Regret Protocol

The selection-regret analysis uses no task-success labels. It evaluates the unaugmented and $\sigma_{\max} = 0.08$ checkpoints on 100 histories per task, training run, and training condition at perturbation severities 0.02, 0.05, 0.08. Each CEM branch uses $K = 64$ candidates, eight elites, eight iterations, and a five-step model-prediction horizon. The clean and perturbed branches start from the same proposal and use common random numbers, but each

Table 9: Transfer of the Gaussian-calibrated selective score to blur and resize. For each task, thresholds are selected on the other three Gaussian-noise tasks and fixed; each pair compares the $\sigma_{\max} = 0.08$ checkpoint with its unaugmented counterpart. “Discordant” counts pairs whose score-change sign disagrees with the observed success-rate label (Section 4.8).

Visual shift	Pairs	BA	Precision / recall	Spearman	Discordant
All pairs	24	0.889	0.882 / 1.000	0.835	2
Blur	12	0.875	0.889 / 1.000	0.873	1
Resize	12	0.900	0.875 / 1.000	0.746	1

Table 10: All 24 LeWM checkpoint pairs evaluated under blur and resize. ATR_{rel} and SMPR are measured for the checkpoint trained with Gaussian augmentation ($\sigma_{\max} = 0.08$); ΔP and ΔS are its success-rate and selective-score changes relative to the unaugmented checkpoint. A pair is positive when $\Delta P \geq 5$ percentage points with at most a five-point clean loss. Daggers mark the two discordant pairs.

Task	Seed	Shift	ATR_{rel}	SMPR	ΔP	ΔS	Outcome (success / score)
TwoRoom	3072	blur	0.499	0.885	55.7	1.670	positive / positive
TwoRoom	3072	resize	0.336	0.967	52.3	2.214	positive / positive
TwoRoom	3073	blur	0.781	0.262	36.0	0.729	positive / positive
TwoRoom	3073	resize	0.691	0.361	40.0	1.030	positive / positive
TwoRoom	3074	blur	0.636	0.541	37.7	1.212	positive / positive
TwoRoom	3074	resize	0.617	0.656	34.3	1.277	positive / positive
PushT	3072	blur	0.943	0.939	10.7	0.189	positive / positive
PushT	3072	resize	0.814	0.969	4.0	0.620	nonpositive / positive [†]
PushT	3073	blur	0.846	0.918	7.3	0.515	positive / positive
PushT	3073	resize	0.682	0.969	24.3	1.060	positive / positive
PushT	3074	blur	0.781	0.959	4.0	0.729	nonpositive / positive [†]
PushT	3074	resize	1.173	0.939	-19.7	-0.577	nonpositive / nonpositive
Reacher	3072	blur	0.155	0.990	50.7	2.375	positive / positive
Reacher	3072	resize	0.210	0.990	29.7	2.375	positive / positive
Reacher	3073	blur	0.176	0.990	52.3	2.375	positive / positive
Reacher	3073	resize	0.207	0.990	40.0	2.375	positive / positive
Reacher	3074	blur	0.190	0.990	44.7	2.375	positive / positive
Reacher	3074	resize	0.195	0.990	33.3	2.375	positive / positive
Cube	3072	blur	1.447	0.330	-2.7	-1.488	nonpositive / nonpositive
Cube	3072	resize	1.166	0.530	1.0	-0.552	nonpositive / nonpositive
Cube	3073	blur	1.577	0.260	2.3	-1.923	nonpositive / nonpositive
Cube	3073	resize	1.218	0.560	-1.3	-0.728	nonpositive / nonpositive
Cube	3074	blur	1.220	0.660	-3.0	-0.735	nonpositive / nonpositive
Cube	3074	resize	1.040	0.770	-2.3	-0.134	nonpositive / nonpositive

branch updates its proposal from its own elites. The implemented cost is the mean-reduced squared error over embedding coordinates, i.e., $\|x - g\|_2^2/d$ for embedding dimension d ; this common positive factor rescales every C_j and b_j identically and leaves the regret bound and both certificates unchanged.

For every candidate, we compute ACPC along that candidate’s action sequence and summarize the pool by its q90 at horizons one and five. For clean plan $\mathbf{a}$ and perturbed-history plan $\tilde{\mathbf{a}}$, the prediction target is the clean-reference selection regret

$$[C_h(\tilde{\mathbf{a}}) - C_h(\mathbf{a})]_+.$$

Within each training run, ridge regressions are fitted on three tasks and evaluated on the fourth. MAE is computed on the log-transformed target and averaged equally across the four tasks excluded from fitting in turn. The reported relative MAE decrease is computed within each run before taking the mean and sample standard deviation across the three runs. This protocol evaluates prediction of a clean-model cost; it does not evaluate simulator return.




Figure 7: Local-sensitivity ratio between the Gaussian-augmented and unaugmented checkpoints, estimated by finite differences and a JVP-based Hutchinson estimator of the composed Jacobian’s squared Frobenius norm. For every task, both ratios are below one, indicating lower local sensitivity after augmentation but not establishing task performance.